

Abstracting Distributed, Time-Sensitive Applications

Kyle Liang

November 2023

School of Computer Science
Carnegie Mellon University
Pittsburgh, PA 15213

Thesis Committee:

Jonathan Aldrich, Chair
Carlee Joe-Wong
Joshua Sunshine
Aviral Shrivastava

*Submitted in partial fulfillment of the requirements
for the degree of Doctor of Philosophy.*

Abstract

Distributed and Time-sensitive (DT) cyber-physical systems are challenging to design and develop. When writing systems in conventional languages, programmers struggle to manage the conciseness and complexity of intrusive, cross-cutting concerns caused by a heterogeneous, distributed system with sensing and actuation timing requirements.

This thesis proposes that a system using the dataflow graph as an intermediate representation can reduce barriers-to-entry for programmers and domain experts unfamiliar with designing distributed, time-sensitive applications. We present TTPython: a domain-specific language and runtime system for writing shorter and cleaner code for challenging cyber-physical applications. Its novel timed token-tagged dataflow graph execution model allows programmers to develop applications at a macroprogramming scale while supporting timing specifications such as periodicity and soft deadlines. The programmer uses annotations to guide TTPython in placement of code and system-provided function calls to specify timing requirements while TTPython handles distribution, communication, and coordination between devices. We have evaluated TTPython by comparing it to a best-practice implementation of a 1/10th-scale connected autonomous vehicle application. An in-progress case study on an intercity flooding application will examine how TTPython affects design and development organically. These two case studies will serve as the foundation of a user study in which we ask users to write DT applications for both of these scenarios using either TTPython or classical Python with a message-broker system.

Dataflow graphs have been applied in various contexts in parallelism and distributed computing, but little work has been focused on token-tagged dataflow graphs. This thesis will show how incorporating time in a token-tagged dataflow graph makes it an effective tool for handling distributed, time-sensitive applications. These advances work towards developing key abstractions that make programming of large arrays of sensors and actuators approachable by non-specialist programmers. With this, we can take steps towards creating a standard digital distribution service for developing and deploying DT applications across devices shared by many, such as in a smart city.

Contents

1	Introduction	1
1.1	Thesis Statement	3
1.2	Evaluation	3
1.3	Outline	4
2	Related Works	5
3	TTPython	7
3.1	Why the Dataflow Graph Architecture?	7
3.2	Example TTPython Program	8
3.3	Time and Location Constraints	10
3.3.1	Specifying Location	10
3.3.2	Deadline Constructs	10
3.4	TTPython’s DFG Semantics	12
3.4.1	Firing Rule: Data-Validity	14
3.4.2	Firing Rule: Time-Based Trigger	14
4	Case Studies	17
4.1	Smart Intersection	17
4.1.1	Code Comparison	18
4.1.2	TTPython Best Practices	21
4.2	Intercity Flooding	22
4.2.1	Initial Findings	23
4.2.2	Methodology	23
4.2.3	Expected Results	23
5	Normative User Study	25
5.1	Comparative User Study between TTPython and Classical Python Solutions . . .	25
5.1.1	Methodology	25
5.2	Expected Findings	28
6	Proposed Contributions	29
7	Proposed Timeline	31

Chapter 1

Introduction

Distributed, time-sensitive (DT) applications permeate cyber-physical systems (CPS). Improvements in wireless, embedded, cloud, and networking technologies enable larger, more interconnected applications that measure and control the physical world. For example, autonomous vehicular networks enable signal-free intersections [13] in which vehicles coordinate their trajectories through the intersection to use the space more efficiently and increase throughput. The vehicles must plan these trajectories with millisecond precision to avoid colliding with each other. Large drone swarms must similarly determine their location and plan their path at millisecond or microsecond precision to form spectacular displays like Intel’s swarm at the 2018 Winter Olympics in South Korea.

DT applications are difficult to develop and deploy, especially when programmers attempt to reason about synchronized, real-time actions in large-scale distributed systems. Programmers are interfacing with dozens if not hundreds of devices with varying hardware capabilities over large geographic regions. Devices collaborate by providing spatial and temporal data and provide valuable information about the environment they reside in by virtue of their hardware or location, but their physical separation causes many technical challenges. DT applications need to include unexpected runtime behaviors of the code, such as recovery procedures for mitigating hardware, software, or network failure. We cannot guarantee hard deadlines due to communication failures and delays, so we must write plans to recover and adapt when deadlines are missed. Distribution and timing concerns are also cross-cutting and intrusive. Coordinating between multiple devices at the device level is unwieldy. The programmer needs to explicitly call 3rd-party library communication protocols to interface between devices. This implicit handshake leads to potential architectural bugs when the application evolves, as message formats and types are hidden by opaque communication function calls. In traditional languages, programmers are restricted by the language’s low-level model of time, which is usually a system call to the UNIX timestamp. This forces programmers to encode the common sense-compute-actuation paradigm in CPSs as an infinite for-loop and compare timing differences on the device’s local clock. The application code awkwardly carries time around with data, which is an inherent requirement in DT apps. Writing recovery procedures for timing violations is also difficult as their installation and execution is dependent solely on time: the deadline by which the action must happen by. The separation between time and data in languages designed without DT systems support in mind leads to verbose and convoluted code.

DT applications require abstractions to handle code distribution, communication, and time requirements at scale. Embedded programming solutions offer helpful abstractions at the device level but fail to scale coordination of space and time for DT applications. Fortunately, prior research in programming languages focus on this issue with macroprogramming. This technique provides a global view of the system and homogenizes the devices in it, offering a powerful abstraction when handling distributed hardware and timing concerns. The programmer focuses on designing system-wide behaviors rather than those local to specific devices. These frameworks provide abstractions for selecting sets of devices, efficient in-network aggregation, and interfaces between heterogeneous devices in the system. This dramatically reduces the overhead for programmers to interface between devices, where the programmer writes a single file rather than one for each device and the framework handles the distribution of code.

We are especially interested in expressing timing constraints across the application. Many research macroprogramming frameworks have been developed but fall short in their handling of timing requirements. Kairos ([8]) was the first step in targeting cyber-physical applications at scale with macroprogramming. It offered extra querying constructs as additions to a host language to easily iterate over multiple devices. However, Kairos lacks timing specifications in its coordination language and requires the user to explicitly couple data with time to handle time requirements. This causes cognitive overhead for the programmer to manage time and data separately in a context where the two in isolation is unhelpful. A modern framework Lingua Franca (LF) ([16]) incorporates a tag with data to specify an ordering of actions in logical time. However, the reduction of timestamps to logical time means LF cannot support the fusion of variable-frequency data streams. LF also requires programmers to check within the host language whether deadlines occur during reaction execution, undercutting the advantages of offering deadline timing abstractions in the macroprogramming layer. In general, macroprogramming research frameworks have little support of timing specifications at the system-level design view. Timing is still handled at the device or function level, whereas DT applications have timing requirements at the global level across devices.

In this thesis, we describe and present TTPython, a domain-specific language and runtime embedded in Python, whose design goals are to help programmers write shorter and cleaner DT applications. TTPython provides abstractions for managing large-scale distributed device networks and for expressing timing constraints on sensing and actuation. Programmers can specify within the programming language soft deadlines in the global distributed application space at a macro level. Our system is written in Python as it is a familiar language for many programmers that is widely supported and has many built-in libraries and useful third-party packages. Many DT applications operate in hardware-constrained environments and have limited power and computation resources, so TTPython does not have a large footprint that could affect system performance because . It also includes constructs for timed, reactive backup plans to account for unreliable hardware and networks.

To integrate macroprogramming into a conventional host language, we compile code to a dataflow graph architecture, which is composed of nodes of computation and edges for data communication. We chose a dataflow graph as it is more flexible in partitioning and in execution timing than its control flow counterpart. The intermediate graph representation supports our systems-level view of the application. This model allows us to easily abstract and assign computation to devices. TTPython utilizes a modified version of the MIT Tagged-Token Dataflow graph

Architecture (TTDA) [20] to represent timing requirements at the macroprogramming level. The TTDA was insufficient to represent time requirements, so we upgraded the architecture by pairing timing information with data and including a control plane to support periodic deadlines. Data is encapsulated in tokens which have a tag to indicate the liveness of the value. The tag includes a time interval to indicate which iteration of periodic execution the data belongs to. The extra control plane added to the TTDA allows us to encode periodic and deadline behaviors at the graph level; thus, the programmer can specify timing at the global system view. To our understanding, we are the first to explore a time-based dataflow architecture.

TTPython removes the overhead of managing multiple devices across separate compilation and the intrusiveness of handling communication between devices. Programmers do not have to worry about using architectural interfaces (be they 3rd party libraries or Internet protocols) to communicate across different devices. Communication is simply a variable reference between different graph nodes (that we call Scheduling Quanta [SQ]) at the system-level view of the program. TTPython focuses on time as a first-class construct. Programmers specify timing requirements and safety mechanisms to take if the timing is not satisfactorily met. Time constraints are specified as data-generation events and deadlines in the global specification of the program. Timing is tightly coupled with data at the system-level view of the application, which highlights timing concerns as a primary concern for the programmer. These constructs are tightly coupled with reactive backup plans, as breaking soft deadlines prompts programmer provided intervention. The programmer seamlessly transitions between developing functional components of their code with Python semantics with the function annotations `@SQify` and `@STREAMify` while specifying their data interactions within the annotated `@GRAPHify` function. TTPython allows non-professional programmers to design and write difficult DT applications.

1.1 Thesis Statement

The design of a macroprogramming language with a modified dataflow graph intermediate representation helps users write safer distributed, time-sensitive applications.

To evaluate this, we examine TTPython through two case studies and a user study.

1.2 Evaluation

We apply TTPython on two realistic distributed, time-sensitive applications: a 1/10th-scale intersection with connected autonomous vehicles (CAV) and intercity flooding sensing. These applications vary in their time and distribution requirements. In terms of development, the intersection application was initially written in classical Python and rewritten in TTPython, whereas the intercity flooding application is using TTPython during its development cycle. These two case studies can shed light on how TTPython affects design and development and the challenges found in using TTPython.

The case studies are the cornerstones of a user study focusing on programmers unfamiliar with cyber-physical systems writing these DT applications. The purpose is to highlight the difficulties programmers face when using conventional frameworks to write DT applications.

Participants will experience writing these applications with both TTPython and classical Python to identify the advantages gained for free with the dataflow graph intermediate representation and the disadvantages of having different semantics from Python as a host language. Each application will have basic tasks for users to complete within a few hours, focusing on making incremental changes to the program. We plan to include a simulation for each application for participants to use to emulate the design and deployment process for these applications.

We will assign programming tasks that are small enough to perform in a few hours but are otherwise representative of the kinds of applications and programming constraints that we observed in our field studies. We will collect data on outcomes, which include time spent programming, bugs introduced (including timing and robustness errors), and success in writing working programs.

1.3 Outline

NOTE: Chapter 2 describes more in-depth description of research that inspired TTPython. We then provide a high-level description of the TTPython system along with its syntax and dataflow graph semantics in Section 3 and Section 3.3.2. Section 4.1 describes a case-study in which we evaluate the experience of designing a non-trivial DT application of a 1/10th-scale intersection with connected autonomous vehicles (CAV). The future-work Intercity Flooding Application is covered in Section 4.2. The work culminates with a comparative user study in Chapter 5.1 to observe the challenges user may face when using TTPython or with other classical solutions.

Chapter 2

Related Works

Wireless Sensor Networks (WSN) [22] are often viewed as a precursor to modern cyber-physical systems, particularly those that rely on wireless communications. WSN research considered a wide variety of programming paradigms [18]. Node-centric programs may be written in low-level languages like C, in domain-specific variants like NesC[7], or as a set of high-level interpreted instructions like Mat  [15]. Popular embedded system languages and frameworks such as PRET-C [1] allow programmers to carefully manage resources on a device-level basis for small applications, but the mechanisms they provide prove unscalable when applications utilize thousands to millions of cost-efficient, resource-constrained devices. The underlying problem is that they lack the abstractions necessary to coordinate CPSs to work at larger scales. Many approaches sought a more holistic view than node-centric programming by introducing abstractions to handle locality via region formation [19, 21], efficient in-network aggregation [9, 17], and shared-memory based on locality [8].

More CPS-focused approaches are cognizant of heterogeneity in the system and the need to clearly define interfaces between elements of the program [3, 4], leading to macroprogramming frameworks. They provide abstractions for selecting sets of devices, efficient in-network aggregation, and interfaces between heterogeneous devices in the system. These often take a host-coordination language approach in which the host language, such as C or Java, is used for platform specific code and the coordination language encodes communication channels, message formats, and which hardware the platform-specific code should run on. These models almost exclusively use message-passing architectures and encode the macroprogram in a graph-based intermediate representation before mapping chunks of code to devices. The devices typically host boilerplate code to handle communication, interfaces, and other common runtime mechanisms that are non-specific to the application. Their functionality reflect many ideas found in Links, which pioneered multitier/tierless programming [5]. By eliminating the interfaces programmers use between different components, Links removed interaction mismatches and centered focus around programmer application code.

Macroprogramming frameworks have slowly begun to include time as a central component. Early frameworks such as Kairos [8] did not include timing constructs within the framework and opted to leave timing handled by its host language, leading to intrusive intermingling of application and system behaviors. COSMOS [3] introduced timing at the system-level overview, but lacks time-specification outside of periodicity. Lingua Franca [16] and PTIDES [23] use "logical

time” in each node for deciding when to act upon time-sensitive inputs. They both introduce timing constructs to specify periodicity and deadlines at the level of devices or functions. This lower granularity of time specification makes it difficult to understand system-wide time requirements. Furthermore, the choice to use logical time at the application level reduces the flexibility of the language when handling variable-frequency data streams.

Importantly, distributed systems interacting with the physical world require a timeline not only for ordering events but also coordinating physical I/O and determining concurrency to compare or combine data values. Having a strong time integration with a macroprogramming framework can alleviate the difficulties found in writing DT applications. The TTPython framework aims to enable easier development of nontrivial distributed, time-sensitive applications by combining these ideas of macroprogramming, time-cognizant architectures, and uncertainty-aware timestamping. The combination of these challenges make this system design difficult, and our advancements in compilation techniques and representation enable TTPython to abstract cross-cutting concerns from distribution while incorporating time-sensitivity.

Dataflow graphs were first introduced in the seminal paper by Dennis [6]. The novelty was to expose the most amount of parallelism in the system by its minimal constraint: data. The semantics were later refined with the U-Interpreter paper [2] by assigning labels to data, allowing nodes to work asynchronously from each other. These ideas were later refined in the MIT Token-Tagged Dataflow paper [20] to work well within the context of a von Neuman machine. Our work builds on the foundation of token-tagged dataflow and introduces time as a specific ”tag”, which allows us to extend dataflow to account for DT applications.

Chapter 3

TTPython

TTPython is a domain-specific language embedded in Python. It separates coordination logic from the application code and allows programmers to specify which code requires careful timing consideration. Although TTPython shares the syntax of Python, top-level application functions execute using a dataflow semantics as specified in Section 3.3.2. The units of computation handled by the TTPython architecture execute in a functional manner such that program flow depends on data availability as opposed to the strict line-by-line imperative execution in Python. However, lower-level functions execute with ordinary Python semantics that is familiar to programmers. TTPython is similar to the Tagged-Token Dataflow Architecture [10, 20] (TTDA) with some modifications to accommodate time, which are detailed in Section 3.3.2. We first motivate the use of dataflow graph architecture before diving into an application to demonstrate TTPython’s capabilities.

3.1 Why the Dataflow Graph Architecture?

When designing a language and system for DT applications, we first take into account the scale of our applications. The number of devices in these applications can be very large, such as in remote sensing where multiple sensors are scattered across a wide area. Coordinating data and timing requirements between all these devices can quickly become unscalable due to complexity in management. On top of this, platform diversity can hinder development as programmers interface with varying frameworks. Thus, the primary decision we start with is for the programmer to write a single, device-agnostic program with abstracted communication and timing requirements. The difficulty now is to partition the program and map the division to the devices.

We utilize the dataflow graph (DFG) as described in the MIT Tagged Token Dataflow Architecture for maximum flexibility during partitioning the program, since data is the minimum requirement for running a block of code. Compiling to an intermediate graph representation also offers great abstractions for cross-cutting concerns such as communication. The dataflow graph matches the macroprogramming approach of declaring system-wide interactions instead of local steps at the device level.

3.2 Example TTPython Program

We describe the language and its semantics through a sample sensor fusion application as seen in Figure 3.1: adding two simulated sine wave data sources. TTPython tracks the addition of these two streams as well as the average of the addition for a period of 30 seconds. The programmer first writes the high-level coordination between functions that will run in parallel in the application. This is done with the function annotation `@GRAPHify`, as seen on line 21. `@GRAPHify` takes the associated function and constructs a dataflow graph from the operations within it. It requires that all operations within the function must either be basic arithmetic or boolean operations or a TTPython annotated function (`@SQify` or `@STREAMify`). The compiler automatically takes basic Python arithmetic and boolean operators and converts them to their SQ counterpart. As the programmer describes the application at the system level, SQs can be assigned to different devices. TTPython uses message passing semantics, so the data passed between SQs are copied and are not shared in memory.

The GRAPHified function `add_sine` is populated by SQs, which are created with the function annotations `@SQify` and `@STREAMify`. `@SQify` takes the associated function and translates it into a node of computation in the dataflow graph intermediate representation. Line 11 shows `@SQify` applied to a moving average function. It weights the existing average by the number of samples, adds the new sample value, and divides by the new number of samples. The implementation of a moving average requires tracking the number of samples taken between iterations. To account for this, TTPython reserves the global variable `sq_state` to store state between multiple executions of a SQ. Note that this does not share the same semantics as Python’s `global` keyword. Instead, the variable is only locally observable by the SQ, so accesses are to the local device’s version of the SQ’s state.

The `@SQify` and `@GRAPHify` constructs we have described so far are insufficient for writing periodic computations. The system needs programmer-specified arguments to describe how long to run the data generation for the sine wave stream. The function annotation `@STREAMify` describes a function that will generate a continuous data stream when called. `@STREAMify` operates similarly to `@SQify` but has added functionality to support periodic execution of the associated function. The function operates similarly to a self-emitting stream in which computation will periodically trigger itself to run. Under the hood, TTPython wraps and unwraps data from arguments and return values respectively through *tokens*, which are values passed along edges in the dataflow graph. Tokens include time intervals with data. Intuitively, these time intervals indicate the time period when the data is valid. Computation involving multiple data values can execute if these data share overlapping time intervals, indicating that they share a common temporal context. We modify an input token’s time interval with the two functions `VALUES_TO_TTIME` and `COPY_TTIME`. `VALUES_TO_TTIME` sets the start and stop ticks of the output token’s timestamp interval to be the values from the two input tokens. `COPY_TTIME` creates a new token with the value of one and the time interval of the other. In Figure 3.1, by setting `N=30`, we are creating a sampling interval of 30 seconds given that the periodicity is 1 second. These are some of the API functions¹ we’ve provided to the programmer to interface with the TTPython architecture.

¹To see a complete view of all the functions we provide, see the supplementary material.

```

1 @STREAMify
2 def sinusoid_sampler(A, f, phi):
3     from math import sin, pi    # local import
4     global sq_state
5     if sq_state.get('count', None) == None:
6         sq_state['count'] = 1
7     sample = A * sin(sq_state['count'] * 2 * f / pi + phi)
8     sq_state['count'] += 1
9     return sample
10
11 @SQify # user defined SQ
12 def moving_average(new_input):
13     global sq_state             # persistent local state
14     count = sq_state.get('count', 0)
15     avg = (sq_state.get('avg', 0) * count + new_input) \
16         / (count + 1)
17     sq_state['count'] = count + 1
18     sq_state['avg'] = avg
19     return avg
20
21 @GRAPHify # main program specifying SQ linking
22 def add_sine(trigger):
23     A_1 = 1; f_1 = 0.25; phi_1 = 0
24     A_2 = 2; f_2 = 0.25; phi_2 = 0
25
26     with TTClock.root() as root_clock:
27         start_time = READ_TTCLOCK(trigger,
28                                     TTClock=root_clock)
29
30         N = 30
31         stop_time = start_time + (1000000 * N)
32
33         # create a sampling interval by setting the
34         # start and stop tick for one of the args
35         sampling_time = VALUES_TO_TTIME(start_time,
36                                           stop_time)
37         A1_sample = COPY_TTIME(A_1, sampling_time)
38         A2_sample = COPY_TTIME(A_2, sampling_time)
39         sine_1 = sinusoid_sampler( # streamify call 1
40                                   A1_sample, f_1, phi_1,
41                                   TTClock=root_clock, TTPeriod=500000,
42                                   TTPhase=0, TTDataIntervalWidth=100000)
43         sine_2 = sinusoid_sampler( # streamify call 2
44                                   A2_sample, f_2, phi_2,
45                                   TTClock=root_clock, TTPeriod=500000,
46                                   TTPhase=0, TTDataIntervalWidth=100000)
47
48         output = sine_1 + sine_2
49         y = moving_average(output)
50     return output

```

Figure 3.1: A TTPython application adding two sine waves and tracking the average periodically every 0.5 seconds for 30 seconds.

The invocation of the STREAMified `sinusoid_sampler` function provides periodicity information via specially named keyword arguments `TTClock`, `TTPeriod`, `TTPhase`, and `TTDataIntervalWidth`. The programmer can specify when the SQ will trigger. A `TTPhase=0` specifies that the `sinusoid_sampler` will trigger when the counter reads 0 modulo the period; for `TTPeriod=500000`, a `TTPhase=250000` would trigger 250ms after the counter wraps around. When running periodic computation, the function also needs to define the time context of the data it generates. When the runtime initiates an instance of the SQ, the runtime will take the start and stop time for executing this SQ and take the average. The new interval is this timestamp average, plus-minus the `TTDataIntervalWidth` divided by 2. The width of the resulting interval is thus `TTDataIntervalWidth`. Currently, this is dictated by the programmer but could later be extended to dynamically estimate the typical clock uncertainty in the network. These special keyword arguments that start with `TT` are specifically used by the runtime and not accessible by the function body during execution.

3.3 Time and Location Constraints

TTPython also offers two more constructs at the `@GRAPHify` level to assist programmers in dealing with space and time constraints. We upgrade our `sinusoid` program to include a deadline constraint on the addition of the two sine wave values and specifying the location of that SQ to a particular device.

3.3.1 Specifying Location

Before graph execution, TTPython needs to decide how to map each SQ in the dataflow graph to its host device. We offer a simple constraint syntax with the keyword `TTConstraint` where the programmer can specify requirements on groups of SQs. These requirements can range from required hardware for an SQ (*e.g.*, a function that takes an image requires a camera on the device) or software (*e.g.*, a specific library for Fast Fourier Transforms). In this program, the programmer has specified that the SQ within the `with` block on line 20 will be assigned to the device named “dev1”. If multiple devices fulfill a requirement, an arbitrary device is selected. Future work will focus on optimizing this selection with respect to timing and other user-specified parameters.

3.3.2 Deadline Constructs

Timing requirements in TTPython have two components: the time by which the computation should finish and the backup procedure to run if the execution doesn’t finish on time. The former is produced by generating a timestamp from `READ_TTCLOCK` and offsetting it with an addition. `READ_TTCLOCK` accepts two parameters: a trigger as to when to take the timestamp and a clock from which to take the timestamp. TTPython uses a clock to allow programmers to specify the level of synchronization necessary between time-sensitive actions. The root clock, a proxy for Universal Coordinated Time, is sufficient for the contributions of this paper. In the updated example in Figure 3.2, `deadline_time` is generated and sent downstream by each iteration of `sinusoid_sampler`. The TTPython SQ `TTFinishByOtherwise` construct ensures that

```

1 @SQify
2 def deadline_check():
3     return None
4
5 @SQify
6 def moving_average(new_input):
7     ...
8
9 @GRAPHify
10 def add_sine(trigger):
11     ...
12     with TTClock.root() as root_clock:
13         ...
14         deadline_time = READ_TTCLOCK(    # deadline generation
15             sine_1,
16             TTClock=root_clock
17         ) + 50000
18         ...
19
20         with TTConstraint(name="dev1"): # assign SQs to 'dev1'
21             safe_add = TTFinishByOtherwise(
22                 output,
23                 TTTimeDeadline=deadline_time,
24                 TTPlanB=deadline_check(),
25                 TTWillContinue=True
26             )
27             # if deadline fails, it produces a separate
28             # value (similar to ternary operation:
29             # a = y if clock.now < t else None)
30
31             moving_average(safe_add)    # downstream SQ

```

Figure 3.2: An updated example of the deadline construct with location specification with adding two sine waves.

data has been generated on time by triggering a backup procedure (*Plan B* [[14]]) if the data is not received by the specified deadline, which is `deadline_check` in our example. If Plan B is run, the programmer can specify different behaviors regarding the rest of execution. These behaviors are similar to those associated with raising an exception.

The SQ `TTFinishByOtherwise` supports two types of recovery behaviors: replacing a “late” value with a default value provided by Plan B or running a separate backup routine without executing the rest of the data dependencies of that data variable. For example, a backup routine on an autonomous vehicle that instructs the car to apply the brakes may not want downstream SQs to execute after applying the brakes, as they could command the car to continue moving forward. `TTWillContinue` accepts a bool and indicates if the programmer intends to execute the downstream SQs from the data value given if Plan B is executed. If `TTWillContinue` is **True**, then the default value passed on will be the value returned by the function call specified with `TTPlanB`. Downstream SQ execution will continue with this de-

fault value. This pass-through procedure is visually represented in Fig 3.3 by the red dotted line. Otherwise, if `TTWillContinue` is **False**, the compiler will not include the red dotted line. The bool value **True** for `TTWillContinue` indicates that the graph will always execute the `moving_average` SQ on line 31 regardless if **Plan B** occurred. The graph will use the value **None** from the output of `deadline_check` as the value for `safe_add` if `sine_1` does not arrive by `deadline_time`.

3.4 TTPython's DFG Semantics

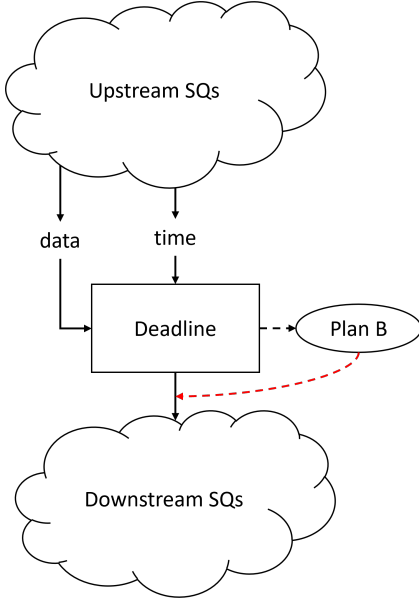


Figure 3.3: Deadline DFG Compilation

We first introduce the original MIT TTDA architecture briefly. The TTDA is a graph that is composed of nodes of computation and edges as communication links between nodes. Data travels through the edges in a form of a token, which encapsulates the actual data with a context. The MIT TTDA is designed to support parallel architectures and can have functions execute simultaneously. The context is used to identify which data belong to the same execution. In TTPython, this context is a time interval. To reflect the emphasis of time in TTPython, we use the term Scheduling Quantum (SQ) for nodes in the dataflow graph. This accounts for the change to a time-specific context for tokens and how the firing rule interacts with this change. The term SQ will be used in the context of TTPython for this section.

A SQ *fires* (executes computation) when certain conditions are satisfied, which is known as its **firing rule**. TTPython has two firing rules: the Data-Validity (DV) and Time-Based Trigger (TBT). The former is the default firing rule for most SQs, as it groups received tokens that match in context to fire. The latter enables the

DFG to encode deadlines at the system level, relieving programmers of low-level code required to interface with time checking. If the firing rule is successful, the SQ then applies its functional computation on the data and updates the context of the token for downstream SQ forwarding. Tokens are stored in a *waiting-matching section*, where tokens wait until the SQ accumulates enough tokens with the same context to fire.

The dataflow graph \mathbb{G} is represented by a set of scheduling quanta V . Each SQ has a set of named input ports and a set of output ports. Edges are implicitly defined by the output ports, which are represented as a list of the names \bar{x} of the input ports they are connected to.

A SQ ($v \in V$) is defined as $\langle f, \bar{i}, \bar{o}, r, s \rangle$ where f is the function v is encapsulating, \bar{i} is a list of input ports, \bar{o} is a list of output ports, r describes the firing rule for v , and s is the internal state of the SQ.

f is of type $\text{List}\langle \text{Token} \rangle \rightarrow \text{Set}\langle \text{Pair}\langle \text{Token}, \text{Name} \rangle \rangle$, where the name is the name of input ports we are sending the corresponding token to (we assume the implementation of the function

does this by looking up what the SQ's output ports are connected to).

Each input port has a waiting-matching section \mathbb{W} , which holds tokens received by the port until the SQ fires. The structure of an input port i is composed of $\langle \mathbf{x}, \mathbb{W} \rangle$. s is the internal state of the SQ for firing rule purposes and differs from the functional `sq_state` as described in Section 3. This is used to keep track of particular tokens for the Time-Based Trigger firing rule, which we will discuss later in Section 3.4.2.

Tokens

An execution of the graph involves tokens propagating through \mathbb{G} . A token k is defined as follows:

$$k = \langle d, t \rangle$$

where d is data and t is a time interval $[t_s, t_e]$ where $t_s \leq t_e$. For simplicity, both timestamps in t can be represented as elements of \mathbb{N} (e.g. a Unix timestamp).

By default, an output token's time interval is the intersection of all the input tokens' time intervals. If the programmer needs to override this default, such as extending the lifetime of a value beyond that of the data it was computed from, we provide APIs for this purpose. These APIs are also used to set the time interval for tokens that specify the length of time to generate a periodic stream, as was shown in Figure 3.1.

DFG Model

We first describe the system configuration \mathcal{C} of DFG. This is described as the tuple

$$\mathcal{C} = \langle \mathbb{G}, \mathcal{N}, t \rangle$$

where \mathbb{G} is the dataflow graph, \mathcal{N} is a set to hold onto tokens in transit through the network, and t as the current time for the system. An element of \mathcal{N} is described as a tuple $\langle k, \mathbf{x} \rangle$ where k is a token and \mathbf{x} is the name of the input port it is addressed to.

Sometimes, there are no possible actions for \mathbb{G} because tokens are still in the network \mathcal{N} . Similarly, tokens can take time to travel through \mathcal{N} . To simulate this, a rule describes the passage of time with no actions taken by the dataflow graph or network. We represent this with the following rule²:

$$\overline{\langle \mathbb{G}, \mathcal{N}, t \rangle} \longrightarrow \langle \mathbb{G}, \mathcal{N}, t + 1 \rangle \quad \textit{Inc-Time} \quad (3.1)$$

Tokens can also exit from the network. They are added to the waiting-matching section of its destination input port.

$$\frac{\begin{array}{l} \langle k, \mathbf{x} \rangle \in \mathcal{N} \quad v \in \mathbb{G} \wedge i \in v.\bar{i} \wedge \mathbf{x} = i.\mathbf{x} \\ v' = v[i \rightarrow \langle x, i, i.\mathbb{W} \cup \{k\} \rangle] \end{array}}{\langle \mathbb{G}, \mathcal{N}, t \rangle \longrightarrow \langle \mathbb{G} \setminus \{v\} \cup \{v'\}, \mathcal{N} \setminus \{ \langle k, \mathbf{x} \rangle \}, t \rangle} \quad \textit{Rcv-Tok} \quad (3.2)$$

The notation for line 2 in Rule 3.2 is shorthand for updating the SQ's input port with the new waiting-matching section unioned with the token to be added. To start a nontrivial execution, an

²This rule has lower precedence than the future rules described later in this section.

initial configuration begins with a nonempty \mathcal{N} . The system generates new tokens by consuming tokens from its input ports and sending them to the network. An SQ's firing rule describes this behavior.

3.4.1 Firing Rule: Data-Validity

An SQ v with the data-validity firing rule will fire if there exists a token on each input port of v such that their time intervals overlap. Tokens $k_i = [t_s^i, t_e^i]$ and $k_j = [t_s^j, t_e^j]$ have overlapping time intervals ($k_i \cap k_j \neq \emptyset$) if $t_s^i \leq t_e^j$ and $t_s^j \leq t_e^i$. The firing rule is described by the rule below:

$$\begin{array}{c}
v = \langle f, \bar{i}, \bar{o}, DV, s \rangle \in \mathbb{G} \wedge \text{len}(\bar{i}) = n \wedge \bar{i} = [i_1, \dots, i_n] \\
\quad \exists k_1, \dots, k_n. \forall a, b \in 1..n. \\
a \neq b \implies k_a \in i_a.W \wedge k_b \in i_b.W \wedge k_a \cap k_b \neq \emptyset \\
\quad f(k_1, \dots, k_n) = \mathcal{N}' \\
\quad \forall j \in 1..n. w_j' = i_j.W \setminus \{k_j\} \\
v' = \langle f, [\langle i_1.x, w_1' \rangle, \dots, \langle i_n.x, w_n' \rangle], \bar{o}, DV, s \rangle \\
\hline
\langle \mathbb{G}, \mathcal{N}, t \rangle \longrightarrow \langle \mathbb{G} \setminus \{v\} \cup \{v'\}, \mathcal{N} \cup \mathcal{N}', t \rangle \quad DV-FR
\end{array} \tag{3.3}$$

3.4.2 Firing Rule: Time-Based Trigger

If a SQ has this firing rule, then it has two input ports, delineated as i_d and i_c . i_d corresponds to a data port, and i_c is the control port. We label tokens k_d and k_c for the i_d and i_c input ports respectively. The list of output ports \bar{o} consists only of downstream SQs that use the data token, so we simplify \bar{o} to a single output port o . The SQ acts as a no-op when the deadline time is satisfactorily met.

The SQ will internally keep track of the last control token it has seen from i_c that caused Plan B to run, which we denote as k_c' . This is tracked in $v.s$, the SQ's internal state. k_c encodes the deadline in its time interval. $k_c.t_s$ denotes when the deadline was initiated and $k_c.t_e$ is the timestamp of the deadline.

TTPython uses a simplifying assumption that all intervals have been widened for uncertainty. Its properties are as follows:

1. Plan B fires at most once per control token. It will either fire if the data token does not arrive on time or be discarded otherwise.
2. If the control token input to a TBT SQ is generated locally, Plan B will run in a timely manner (*i.e.*, there are no network complications in running the SQ).

We describe the rules with success and failure cases. Summarily, success is when both conditions are satisfied:

1. Data is synchronized with the control token (*i.e.*, they overlap).
2. Tokens arrive on time (*i.e.*, they arrive before the deadline specified by the control token).

$$\begin{array}{c}
v = \langle f, [i_d, i_c], o, \text{TT}, \langle k'_c \rangle \rangle \in \mathbb{G} \\
k_d \in i_d.\mathbb{W} \quad k_c \in i_c.\mathbb{W} \quad t \leq k_c.t_e \quad k_d \cap k_c \neq \emptyset \\
\bar{i} = [\langle i_d.\mathbf{x}, i_d.\mathbb{W} \setminus \{k_d\} \rangle, \langle i_c.\mathbf{x}, i_c.\mathbb{W} \setminus \{k_c\} \rangle] \\
v' = \langle f, \bar{i}, o, \text{TT}, \langle k'_c \rangle \rangle \\
\mathcal{N}' = \{\langle k_d, x \rangle \mid x \in o\} \\
\hline
\langle \mathbb{G}, \mathcal{N}, t \rangle \longrightarrow \langle \mathbb{G} \setminus \{v\} \cup \{v'\}, \mathcal{N} \cup \mathcal{N}', t \rangle \quad \text{\textit{TBT-S}}
\end{array} \tag{3.4}$$

Exceptional cases occur when either condition is not met.

- Data is late (*i.e.*, time t is greater than the control token k_c 's end timestamp).

Plan B is run (Rule 3.5), and we will discard any data tokens that would be synchronized to this control token (Rule 3.6). The function $\max(k, k')$ returns the token with the greater t_e .

$$\begin{array}{c}
v = \langle f, [i_d, i_c], o, \text{TT}, \langle k'_c \rangle \rangle \in \mathbb{G} \\
k_c \in i_c.\mathbb{W} \quad k_c.t_e < t \\
i_c^n = \langle i_c.\mathbf{x}, i_c.\mathbb{W} \setminus \{k_c\} \rangle \\
v' = \langle f, [i_d, i_c^n], o, \text{TT}, \langle \max(k_c, k'_c) \rangle \rangle \\
f(k_1, \dots, k_n) = \mathcal{N}' \\
\hline
\langle \mathbb{G}, \mathcal{N}, t \rangle \longrightarrow \langle \mathbb{G} \setminus \{v\} \cup \{v'\}, \mathcal{N} \cup \mathcal{N}', t \rangle \quad \text{\textit{TBT-F-PBC}}
\end{array} \tag{3.5}$$

$$\begin{array}{c}
v = \langle f, [i_d, i_c], o, \text{TT}, \langle k'_c \rangle \rangle \in \mathbb{G} \\
k_d \in i_d.\mathbb{W} \quad k_d.t_s \leq k'_c.t_e \\
i_d^n = \langle i_d.\mathbf{x}, i_d.\mathbb{W} \setminus \{k_d\} \rangle \\
v' = \langle f, [i_d^n, i_c], o, \text{TT}, \langle k'_c \rangle \rangle \\
\hline
\langle \mathbb{G}, \mathcal{N}, t \rangle \longrightarrow \langle \mathbb{G} \setminus \{v\} \cup \{v'\}, \mathcal{N}, t \rangle \quad \text{\textit{TBT-F-PBD}}
\end{array} \tag{3.6}$$

- Data has arrived but is unsynchronized.

This failed condition has three cases:

- **Data token has arrived but there is no overlapping control token.**

The SQ will wait for a control token to come.

- **Control token arrives after its specified deadline.**

As the deadline has already passed, the SQ must run Plan B immediately. This is also already encoded with Rule 3.5.

- **Control and Data token do not match.**

This rule depends on the order of the time intervals between the control and data tokens. We assume here that $k_d \cap k_c = \emptyset$.

If k_d 's time interval is before k_c , the SQ will wait until time is greater than $k_d.t_e + \text{Exp}$, where Exp is a programmer-defined time extension to wait for the control token. Exp by default is upper-bounded by the periodicity of the stream that k_d has been generated by. The default ensures in-order operation for periodic streams. This implies

that if iteration n of the control token in the stream arrives, any tokens from iteration $n - 1$ are automatically late (unless if Exp is set explicitly longer by the programmer).

$$\begin{aligned}
v &= \langle f, [i_d, i_c], o, \text{TT}, \langle k'_c \rangle \rangle \in \mathbb{G} \\
&\quad k_d \in i_d.\bar{W} \quad k_c \in i_c.\bar{W} \\
&\quad k_d.t_e < k_c.t_s \quad k_d.t_e + \text{Exp} < t \\
&\quad i_d^n = \langle i_d.x, i_d.W \setminus \{k_d\} \rangle \\
&\quad v' = \langle f, [i_d^n, i_c], o, \text{TT}, \langle k'_c \rangle \rangle \\
\hline
\langle \mathbb{G}, \mathcal{N}, t \rangle &\longrightarrow \langle \mathbb{G} \setminus \{v\} \cup \{v'\}, \mathcal{N}, t \rangle \quad \textit{TBT-F-US}
\end{aligned} \tag{3.7}$$

If k_c 's time interval is before k_d , the SQ waits for the corresponding data and control tokens for each as the data for the next iteration has arrived earlier than expected.

Chapter 4

Case Studies

4.1 Smart Intersection

Cooperative autonomous vehicle intersections present a unique problem in terms of timing. The goal of an autonomous vehicle intersection is to allow the Connected Autonomous Vehicles (CAVs) to drive through the intersection as close to the speed limit as possible without colliding [13]. Often this will result in vehicles traveling at high velocity in close proximity to each other with little margin for error. Accomplishing this feat requires precise timing of sensing, sensor fusion, communications, and execution of the planned path. Any deviation or failure could be catastrophic. This case study application was created using 1/10-scale autonomous vehicle models with scale accurate LIDAR and camera sensors shown in figure 4.1 as well as 1/10-scale connected infrastructure sensor (CIS) which is a stationary camera that is also sharing information with the intersection, shown in figure 4.2. The 1/10-scale vehicles drive a figure eight loop with an intersection in the middle, which can be seen in Figure 4.3. This intersection is controlled using an autonomous vehicle intersection controller that is run on a Jetson TX2. The intersection has four scale CAVs to control, 2 scale CISs, and a road side unit (RSU) that is running the intersection controller.




Figure 4.1: One tenth scale CAV with camera, LIDAR, and Nvidia Jetson Nano for on-board processing.




Figure 4.2: One tenth scale CIS with camera using Nvidia Jetson Nano for on-board processing.




Figure 4.3: Overhead of one tenth scale CAVs shown driving within the figure 8 intersection.

The CAV intersection application consists of a multitude of sensor streams that must be

combined into a single worldview, which is then used to determine the paths each CAV will carry out in the physical world. Figure 4.4 depicts a CAV or CIS data-flow diagram. All processes of the CIS are depicted in blue while those of the CAV are in blue and gray. A CAV has an extra LIDAR sensor and can actuate a steering and drive motor to move, whereas the CIS sensor is just for sensing alone and has no actuation ability. All CAVs and CIS sensors in the network must synchronise their sensor frequency to within a tight margin so that the sensor fusion operates correctly. All CAVs and CISs must process and locally fuse their data and have it sent to the RSU so that there is enough time for the RSU to calculate the global fusion and intersection control [12] (see Figure 4.5). Then that data is in turn sent out to all the CAVs so they actuate their steering and motors on time. This entire process must happen within 0.125ms so that timing error does not cause a crash of the CAVs. This timing and communication is very sensitive to problems and thus makes an excellent case study for TTPython. Even slight timing issues can cause the perceived positions and trajectories of the CAVs to be estimated incorrectly which can result in a crash as the CAV is in actuality at a different position. Additionally any late or out of order communications can result in a crash because these type of errors may also affect the perceived positions.




Figure 4.4: Data-flow of a 1/10-scale Connected Autonomous Vehicle (CAV) shown in blue and gray and a 1/10-scale Connected Infrastructure Sensor (CIS) shown in blue. CAVs and CISs send their locally fused sensor data to the RSU and CAVs receive intersection control back which is used to actuate the steering and motors.

4.1.1 Code Comparison

We compare the original code for the aforementioned smart intersection application to a newly converted TTPython version. The original application has 6415 lines of code, 624 of which is non-application specific code. The functionality of both code bases is identical from the application output perspective. However, many of the internal structures were changed in migrating the original Python code base to TTPython. These changes fall into three main categories: 1) timing management code ensuring that sensors and devices sense and actuate at the proper time, 2) threading and thread management, including pipelines for thread-to-thread communication, and 3) networking code allowing all the different devices to communicate. TTPython reduces




Figure 4.5: Data-flow of a 1/10-scale Road Side Unit (RSU) that gathers sensing data from the CAVs and CISs in the area, processes the global sensor fusion, and calculates the intersection controls to send back out to the CAVs in the area.

(LoC)	Original App	TTPython Variant
Timing	255	39
Threading	29	X
Networking	340	X

Table 4.1: Breakdown of non-application specific lines of code. TTPython gets threading and networking code for free due to the nature of the DFG.

the code needed for all three of these categories and is able to eliminate 93% of the code in these three categories combined, as seen in Table 4.1.

Timing Management Code is Drastically Reduced

Timing management code manages the flow of a program in time. This includes making sure that information sources such as the cameras and LIDARs stay synchronized at the desired frequency and within a certain time interval. It also includes checking for missing data and running backup routines (PlanB). TTPython provides direct support for timing management, so there is marked improvement in which 255 lines of Python code were removed and replaced by 39 lines of TTPython. For example, in TTPython a single line of code such as `cam_sample = camera_sampler(cav_0, sample_window, TTClock=root_clock, TTPeriod=125000, TTPhase=0, TTDataIntervalWidth=62000)` can replace around 30 lines of manually written Python code to directly manage the timing and synchronization. Thus, as shown in Table 4.1, our case study reduced this category of code by 85%, while also making the code more declarative and easier to understand.

Examining Figure 4.6, the original CAV application made direct use of the Python time library to implement waiting for the right time to start computation and a `while` loop to implement periodicity with a `time.sleep()` to run every millisecond. As seen in Figure 4.7, TTPython separates these concerns from the programmer into 4 steps: setting a start time, stop time, periodicity, and data interval validity for downstream tokens. In the interest of space, we have not shown the timestamp checking code for the original application. This code is built into the TTPython runtime system and thus does not have to be written by the applications programmer; eliminating it is one of the reasons for the reduction in code size mentioned above.

```

# Start after 10 seconds
start_time = 10000000
interval = 125000 # Interval is 125ms

# Sleep until test start time
wait_until_start = (
    start_time - time.time() -.01
)
if wait_until_start > 0:
    time.sleep(wait_until_start)

# Sleep thread until target time
target = start_time
while 1:
    if (fetch_time(simulation_time) >=
        target):
        now = time.time()

        camcoordinates = camera_sampler()

        # Prep value to be sent
        # Clear the queue of old data
        while not out_queue.empty():
            out_queue.get()
            out_queue.put(
                [camcoordinates,
                 now,
                 time.time()]
            )

        target = target + interval
        time.sleep(.001)

```

Figure 4.6: *The original CAV application using Python time and user-written communication library.*

```

N = 10
# starts N seconds later
start_time = READ_TTCLOCK(
    trigger, TTClock=root_clock
) + 1000000 * N
stop_time = GET_INFINITY(
    trigger, TTClock=root_clock
)
sampling_time = VALUES_TO_TTTIME(
    start_time, stop_time
)
sample_window = COPY_TTTIME(
    A_1, sampling_time
)

with TTConstraint(name="cav0"):
    cam_sample = camera_sampler(
        cav_0, sample_window,
        TTClock=root_clock,
        TTPeriod=125000,
        TTPhase=0,
        TTDataIntervalWidth=62000
    )

```

Figure 4.7: *The CAV application rewritten using TTPython.*

Thread Management and Networking Management Code is Eliminated

In TTPython, threading is managed by the data flow runtime. From the programmer’s perspective, each SQ can be considered its own thread, and thus the need to write extra code to manage each thread (beyond `SQify` annotations) is not necessary. This completely eliminates 29 lines of threading and thread management from the original application. Communication is also built into the TTPython data-flow architecture. In our case study the user had to select a communication stack and chose Flask, a Python library. Though a popular choice for networking, it was still not trivial to implement the communication API and call it from each device. Additionally, timing data structures, fallback routines, and routing tables had to be created manually, whereas TTPython generates those automatically. TTPython thus resulted in a savings of 340 lines of code, as well as the complete removal of the Flask library.

Referring back to Fig 4.6, there are explicit calls to the network queues used for connections to other ensembles. In the old application, each ensemble required a separate file responsible for handling communication. TTPython provides this to the programmer for free. In TTPython, communication across devices is handled via variables and thus a line of code like `local_fusion_result = local_fusion(cam_output, lidar_output)` followed by `global_fusion_result = global_fusion(local_fusion_result)` specifies that the local fusion result from up to n CAVs will be sent via TTPython’s networking stack, accumulated and fed to the function `global_fusion()` with no further networking or synchronization code necessary from the programmer side. This further simplifies the project structure as shown in Fig 4.8. In the original application, there are separate folders for the CAV and RSU code. Any shared communication would first need to interface with the networking library before other ensembles could receive it. Sharing is simple in TTPython, as `with` blocks with `TTConstraint` do not limit variable scoping. Referencing and updating what data needs to be shared across ensembles is more explicit and easier to verify within a file as a shared variable name.

4.1.2 TTPython Best Practices

The process of converting to TTPython is trivial at times and at other times required a complete re-arrangement of how a traditional programmer may write code. TTPython effectively requires the programmer to consider the distributed data flow nature of their code at all points. For instance, reading from the camera and LIDAR sensors using the `@STREAMify` construct happens in TTPython simultaneously despite being written in a sequential fashion whereas in traditional Python the functions would be executed in-order. This allows inherent parallel execution by simply setting a timing parameter in `@STREAMify` rather than needing to explicitly split code into threads or use the thread pooling construct. This mindset is very useful, because as is well documented in the litera-

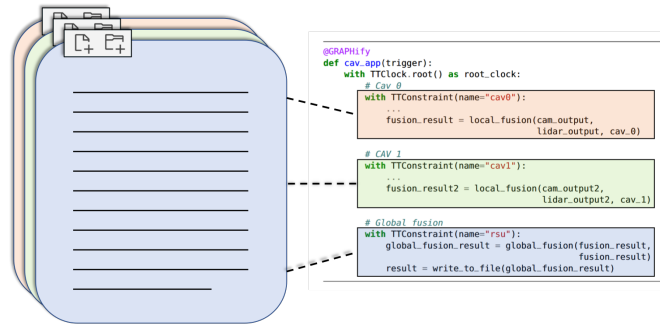


Figure 4.8: Macroprogramming framework TTPython Code. The original implementation had separate files designated for each ensemble, while TTPython supports development of all constituent ensembles seamlessly in one file.

ture, distribution is not a feature that can be hidden but rather should be at the forefront of the programmer’s mind [11]. TTPython is not designed for all distributed systems; for example, if a programmer wants to write a single threaded sequential type of code that connects to a server and uploads some data, TTPython will not be useful for that programmer. However a programmer writing highly distributed and multi-threaded code executing on heterogeneous devices would benefit from TTPython’s abstractions for managing distribution, concurrency, and mapping to devices. In our Intelligent Intersection use case, the programmer gains a lot for free in terms of timing and sequentiality of the firing order of processes as well as many benefits from the communication protocols that are built into TTPython functions.

Built-in Communication, Timing, and Synchronization Saves Time

The second source of improvement from TTPython is that communication, timing, and synchronization are built-in. Thus, the programmer does not need to worry about using libraries or implementing these manually. This is a huge benefit in terms of development-time savings. By simply calling an SQ and using the resulting variable in the next SQ, TTPython will automatically handle that communication, including determination of the downstream code’s location on the network and correctly transferring data to said downstream code, be it on another device or the current one. There is no need to implement APIs for networking or to understand exactly where the SQ is mapped—if it is on the same device, TTPython will use short-circuit the networking stack to save time.

By using the TTPython `@STREAMify` decorator, almost all aspects of timing for a sensor stream can be initialized using a single line of code. In fact, that line of code does even more, incorporating validity intervals that will support semantic matching with other similarly produced tokens as they arrive to in downstream SQs. This saves a significant amount of coding effort: the programmer need not recreate the timing structures manually nor concern themselves with time-stamping data and matching the validity intervals. Finally, the backup routines present in the Plan B structure within TTPython allows for the easy addition of any backup function required with only two extra line of code and no additional timestamp comparisons.

4.2 Intercity Flooding

The second case study will focus on wide-area sensing for detecting flooding of intercity sewer systems. Elizabeth Carter from Syracuse University is working closely with the United States Geological Survey (USGS) and the city of Syracuse to develop this application. Uneven water levels in a city sewage system can be caused by numerous external factors, such as impeding foliage or focused rain fall. The application combines thermal and optical imaging with location sensors (GPS and IMU) to identify where flooding occurs within the sensor network deployed across the city. If flooding is detected, the app will notify personnel to address the flooding. Furthermore, the device that detected the flooding will notify surrounding devices to increase their sampling rate to better gauge the water levels as the situation unfolds. The city of Syracuse plans to integrate this application into its city infrastructure, which can provide a more stable environment in terms of power and internet capabilities, while the USGS is interested in more remote locations where cell reception can be difficult to find.

We are already in an active collaboration with Dr. Carter to author their Intercity Flooding application. The application is an ideal real-world testbed for TTPython, as it is a distributed, cyber-physical system including timing constraints which vary adaptively, failure scenarios that motivate backup routine use, and different modalities of execution based off of environmental factors. We believe that TTPython can expand

the original scope of the Intercity Flooding project due to the abstractions it provides. It is motivated by these research questions.

RQ 1: How do domain-scientists work with experts in cyber-physical systems to create DT applications?

RQ 2: How does timing and distribution appear in the design and development of DT applications?

RQ 3: How do developers interact with timing and distribution requirements in TTPython?

4.2.1 Initial Findings

There have been initial observations with Dr. Carter’s team both at the development and planning level to better understand the scope of the flooding application. These observations were helpful in identifying our research questions. The next step is to formalize this data collection as specified in Section 4.2.2.

4.2.2 Methodology

To collect data for these research questions, we will conduct a series of interviews and observations of the team. These will be semi-structured interviews of the team members in Dr. Carter’s group asking about their experiences before and after using TTPython on the project. The observations will include recorded sessions with the participant using TTPython to write the flooding application. The programmer will use think-aloud to express their thought process when developing the application through TTPython. We will also record interactions with the Dr. Carter’s team with her collaborators that are familiar with cyber-physical systems to understand how Dr. Carter brings her expertise in applied computational hydro-climatology to an application on cyber-physical systems.

4.2.3 Expected Results

For RQ 1, we expect to see that domain-scientists have a strong understanding of the model of the environment where they want to deploy their application. For example, Dr. Carter is an expert in combining imaging and satellite data to identify flooding. However, we have noticed that their understanding of realizing this application onto actual hardware is quite different. Timing and distribution, which are crucial in DT applications, are not their primary concern. We expect to see more evidence of this through the semi-structured interview and through future meeting observations.

RQ 2 will be better answered in the future once the application is further developed. There have been some initial design meetings in which we have discussed timing and distribution, but nothing concrete has been made yet.

We expect to see that at the programming level, TTPython plays a large role in focusing development of code in terms of timing and distribution. TTPython requires programmers to first consider the high-level timing and distribution requirements before writing code while abstracting necessary boilerplate code required to coordinate and parallelize the execution. Evidence will come from the proposed think-aloud recorded programming sessions.

Chapter 5

Normative User Study

5.1 Comparative User Study between TTPython and Classical Python Solutions

To validate TTPython’s efficacy in abstracting timing and distribution concerns in DTS applications, we will conduct a user study with participants developing two sample programs. These two sample programs will have a corresponding solution in TTPython and Python with library support for timing and distribution, namely the Python `time` library and a message broker (RabbitMQ). We believe this to be worthwhile comparison as both case studies use Python and its basic time library to handle their DT application. Furthermore, many popular embedded frameworks use a publish-subscribe model to handle communication between distributed embedded devices.

TTPython’s target audience include domain experts unfamiliar about CPS-specific concerns found in DTS applications. The difficulty in this user study is to introduce participants to non-trivial DTS applications while abstracting most of the domain knowledge required to develop these programs. We expect that TTPython will show advantages both in describing timing and distribution requirements and in software evolution concerning these requirements. This is due to the fact that time is inherently coupled with data in tokens and communication is implicitly specified by variable call across SQs in the dataflow graph. Our study aims to answer the research questions listed below.

- RQ 1: Does macroprogramming reduce the overhead of communication with managing applications in a distributed, time-sensitive environment?**
- RQ 2: What are the challenges programmers face when decomposing a sequential program into one compatible with a dataflow graph architecture?**
- RQ 3: How does each framework handle DTS program evolution?**

5.1.1 Methodology

We will solicit participants from CMU and acquaintances and select experienced Python programmers. Each participant will self-report their level of expertise in Python.

Each participant will use TTPython and classical Python to implement one application. Thus, each participant will first take a tutorial with periodic quiz questions to foster active learning. Participants will

be randomly assigned both the order and which solution to use per application. Both applications contain two parts: the first in which the participant will be handed a finished application without timing or distribution requirements. This scaffolding comes from the idea that the application was designed and tested on a single device before wide-scale distribution, which should ease the burden of understanding domain-specific knowledge of each program. The programmer will then take this single application and appropriately add timing and distribution requirements through TTPython or classical Python, the details of which are specified in later subsections. The second part will have the participant take a solution (either the participants or we will provide ?) and evolve the application due to new timing and distribution requirements. This reflects the nature of the difficulty in writing general DTS applications, as these applications need to be tuned and appropriately scaled to the environment they run in. These applications are difficult to work properly on initial launch, so multiple re-deployments are expected.

To make the development process as realistic as possible, we plan to provide simulated data streams and output monitors for participants to use to run and debug their applications. Both applications will use a simulated data source to provide a deterministic output. For the 1/10-scale CAV Intersection, we will have a 2D-graph plotting the movement of the cars through the intersection. Its data will be sourced from a sample data source provided by our collaborators at Arizona State University. The intercity flooding detection will have a command-line interface printing a log describing the output of the sensors. We will base our simulated data on prior flooding data taken from the USGS monitoring website.

We will be measuring the time taken to complete the program, bugs encountered during development and in the solution, and lines of code for solution. An additional metric to take specifically during the software evolution portion includes the number of lines changed between programs. This is to understand how intrusive timing and location requirements are for each design to accommodate.

Task 1: 1/10th-scale Connected Autonomous Vehicle Intersection

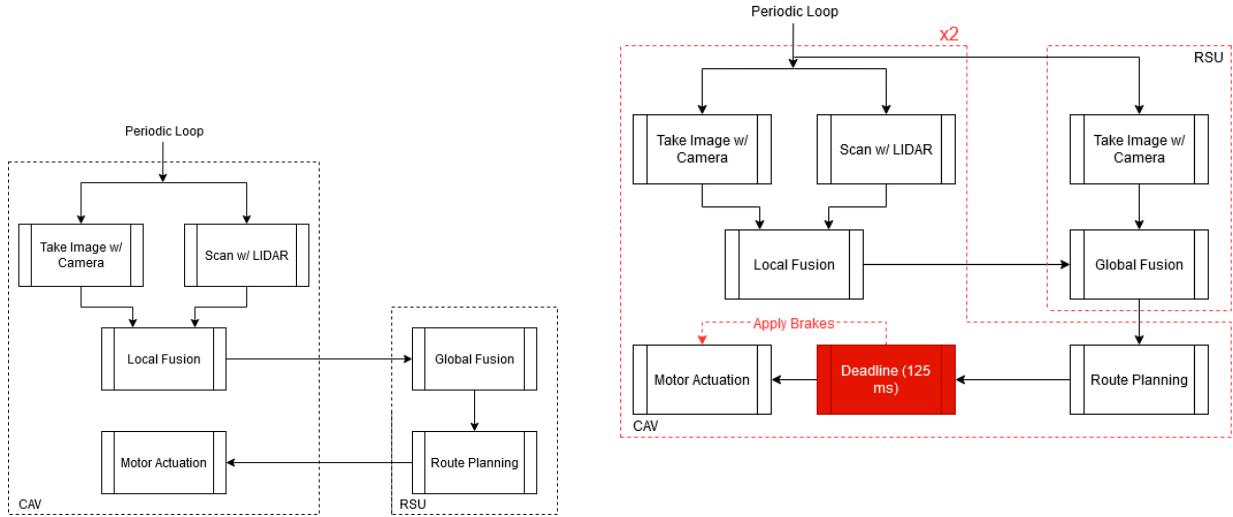


Figure 5.1: A dataflow graph of the phase 1 solution to the CAV application.

Figure 5.2: A dataflow graph of the phase 2 solution to the CAV application. The components highlighted in red show the difference between phase 1.

Participants will write a 1/10-scale Connected Autonomous Vehicle Intersection application. The application consists of three types of devices: connected autonomous vehicles, infrastructure sensors, and

road-side units. Each model autonomous car is equipped with camera and LIDAR sensors and performs local fusion to detect nearby objects. They then communicate with each other combined with an infrastructure sensor for global fusion to collaborate on the objects they detect. Global fusion occurs on the road-side unit, the host for this global state of the environment of the intersection. If there are disruptions in the periodic execution of 125 milliseconds, the cars will take cautionary procedures such as braking.

For the first part, participants will be given this application designed to run for one device: namely an untimed for-loop containing local fusion and global fusion. Participants will have to both set timing requirements (of periodicity and matching-in-time data values) and partition the program to two devices: the car itself and a road-side unit. The dataflow graph of this program is seen in Figure 5.1

Once finished, the requirements will change to a larger-scale deployment, which can be seen in Figure 5.2. The application will use two cars and need a timed backup procedure in case the program fails to meet the deadline requirement of 125 milliseconds. The CAV car associated with the late response will apply the brakes to avoid crashing due to limited knowledge of the surrounding environment. We also have participants move some of the computing from the RSU to the CAVs (particularly the route planning). The code is also modified for the RSU to also take an image of the intersection to represent the infrastructure sensor as is used in the original application.

Task 2: Intercity Flooding Detection

Participants will develop a wide-area sensor application to detect flooding of intercity sewer systems. Uneven water levels in a city sewage system can be caused by numerous external factors, such as impeding foliage or focused rain fall. The application combines thermal and optical imaging with location sensors (GPS and IMU) to identify where flooding occurs within the sensor network deployed across the city. If flooding is detected, the app will notify city personnel to address the flooding.

Participants will initially write a sample program with Figure 5.3 dataflow graph in mind. Users will create three separate data streams: two for imaging and one for coordinate collection. The data streams will be used in a machine learning model as a function call to determine if flooding is detected. If it is, the device will then upload an image and coordinates to a remote server to simulate alerting city officials.

Participants will then upgrade their program to add the components highlighted in red seen in Figure 5.3. This includes adding a Plan B procedure to reset the GPS and IMU if it is taking too long to reach a signal and modifying the periodic sampling time based off of the flooding condition.

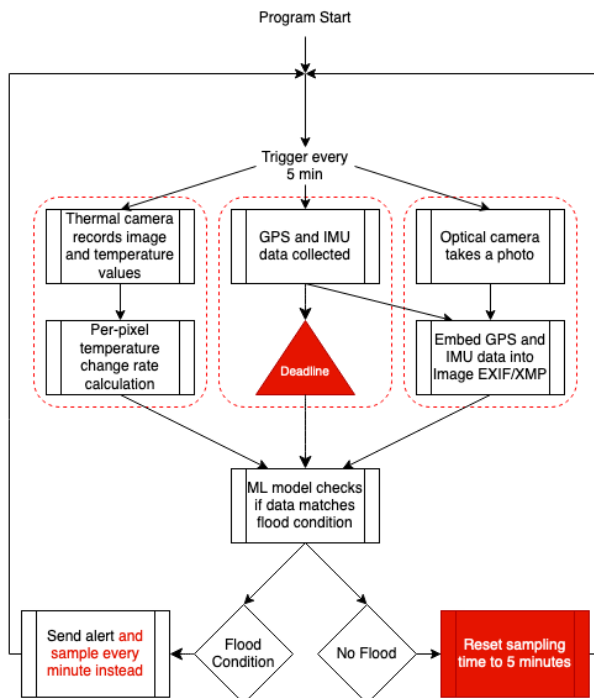


Figure 5.3: A sample dataflow graph of the solution to the Flooding Detection application. The components highlighted in red are to be completed in the 2nd phase of development.

5.2 Expected Findings

We expect to see that TTPython will have fewer errors and faster completion. This is due to the abstractions that TTPython provides with distribution and timing. In DT applications, it is essential that data to be used together in computation share a temporal context (i.e. their timestamps are close together in time). The code to do so is difficult to write and manage, as this must be done consistently throughout the program and is often required when sending data between devices.

We also expect programmers to have challenges with the dataflow semantics in Python. Classical Python is imperative, and our use of macroprogramming changes the logic of some keywords in Python such as `global`. The tutorial may mitigate some of these challenges.

For RQ 3, we expect to see TTPython incorporate changes easily. Due to its macroprogramming nature, communication is abstracted into variable calls and scopes. This is not the case for message passing systems, as it requires to open a new topic or channel. We expect this to be a major overhead when incorporating changes in pre-existing code bases.

Chapter 6

Proposed Contributions

My thesis will have made the following contributions:

- A defined syntax and semantics for a language and framework (TTPython) for an intermediate dataflow graph execution.
- Evidence of advantages of TTPython for abstracting timing and distribution concerns through a case study and user study.
- Evidence showcasing challenges domain-experts face when working in a cyber-physical environment.
- Evidence of how dataflow graphs can manage code evolution in DT applications.

Chapter 7

Proposed Timeline




Figure 7.1: A breakdown of deliverables for the thesis. White is the estimated amount of work finished for the task, while gray is unfinished.

I plan to finish up and submit the TTPython paper to TOPLAS within this year. I am waiting on some new results on the implementation of Plan B. The case study and the user study can be done in parallel as they are not dependent on each other: the user study is simply using the intercity flooding case study as a practical, non-trivial example for a task. I have left some time in the summer on the case I take an internship. I expect that the user study will take 3 months to complete.

Bibliography

- [1] Sidharta Andalam, Partha Roop, Alain Girault, and Claus Traulsen. Pret-c: A new language for programming precision timed architectures. 2009. 2
- [2] A. Arvind and K. Gostelow. The u-interpreter. *Computer*, 15(02):42–49, feb 1982. ISSN 1558-0814. doi: 10.1109/MC.1982.1653940. 2
- [3] Asad Awan, Suresh Jagannathan, and Ananth Grama. Macroprogramming heterogeneous sensor networks using cosmos. *ACM SIGOPS Operating Systems Review*, 41(3):159–172, 2007. 2
- [4] Amol Bakshi, Viktor K Prasanna, Jim Reich, and Daniel Larner. The abstract task graph: A methodology for architecture-independent programming of networked sensor systems. *Proceedings of the 2005 workshop on Endtoend senseandrespond systems applications and services*, (Eesr 05):19–24, 2005. 2
- [5] Ezra Cooper, Sam Lindley, Philip Wadler, and Jeremy Yallop. Links: Web programming without tiers. In Frank S. de Boer, Marcello M. Bonsangue, Susanne Graf, and Willem-Paul de Roever, editors, *Formal Methods for Components and Objects*, pages 266–296, Berlin, Heidelberg, 2007. Springer Berlin Heidelberg. ISBN 978-3-540-74792-5. 2
- [6] Jack B. Dennis and David P. Misunas. A preliminary architecture for a basic data-flow processor. In *Proceedings of the 2nd Annual Symposium on Computer Architecture*, ISCA '75, page 126–132, New York, NY, USA, 1974. Association for Computing Machinery. ISBN 9781450373661. doi: 10.1145/642089.642111. URL <https://doi.org/10.1145/642089.642111>. 2
- [7] David Gay, Philip Levis, Robert Von Behren, Matt Welsh, Eric Brewer, and David Culler. The nesc language: A holistic approach to networked embedded systems. *Acm Sigplan Notices*, 38(5):1–11, 2003. 2
- [8] Ramakrishna Gummadi, Nupur Kothari, Ramesh Govindan, and Todd Millstein. Kairos: a macro-programming system for wireless sensor networks. In *Proceedings of the twentieth ACM symposium on Operating systems principles*, pages 1–2, 2005. 1, 2
- [9] Timothy W Hnat, Tamim I Sookoor, Pieter Hooimeijer, Westley Weimer, and Kamin Whitehouse. Macrolab: a vector-based macroprogramming framework for cyber-physical systems. In *Proceedings of the 6th ACM conference on Embedded network sensor systems*, pages 225–238, 2008. 2
- [10] Robert A Iannucci. Toward a dataflow/von neumann hybrid architecture. *ACM SIGARCH Computer Architecture News*, 16(2):131–140, 1988. 3
- [11] Samuel C. Kendall, Jim Waldo, Ann Wollrath, and Geoff Wyant. A note on distributed computing. Technical report, USA, 1994. 4.1.2
- [12] Mohammad Khayatian, Rachel Dedinsky, Sarthake Choudhary, Mohammadreza Mehrabian, and Aviral Shrivastava. R2im – robust and resilient intersection management of connected autonomous

- vehicles. In *proceedings of The 23rd IEEE International Conference on Intelligent Transportation Systems*, 2020. 4.1
- [13] Mohammad Khayatian, Mohammadreza Mehrabian, Edward Andert, Rachel Dedinsky, Sarthake Choudhary, Yingyan Lou, and Aviral Shrivastava. A survey on intersection management of connected autonomous vehicles. *ACM Transactions on Cyber-Physical Systems*, 4(4):1–27, 2020. 1, 4.1
 - [14] Mohammad Khayatian, Mohammadreza Mehrabian, Edward Andert, Reese Grimsley, Kyle Liang, Yi Hu, Ian McCormack, Carlee Joe-Wong, Jonathan Aldrich, Bob Iannucci, and Aviral Shrivastava. Plan b - design methodology for cyber-physical systems robust to timing failures. *ACM Trans. Cyber-Phys. Syst.*, jan 2022. ISSN 2378-962X. doi: 10.1145/3516449. URL <https://doi.org/10.1145/3516449>. Just Accepted. 3.3.2
 - [15] Philip Levis and David Culler. Maté: A tiny virtual machine for sensor networks. In *Operating Systems Review (ACM)*, volume 36, pages 85–94. ACM, 2002. doi: 10.1145/635508.605407. 2
 - [16] Marten Lohstroh, Christian Menard, Soroush Bateni, and Edward A Lee. Toward a lingua franca for deterministic concurrent systems. *ACM Transactions on Embedded Computing Systems (TECS)*, 20(4):1–27, 2021. 1, 2
 - [17] Samuel R Madden, Michael J Franklin, Joseph M Hellerstein, and Wei Hong. Tinydb: An acquisitional query processing system for sensor networks. *ACM Transactions on database systems (TODS)*, 30(1):122–173, 2005. 2
 - [18] Luca Mottola and Gian Pietro Picco. Programming wireless sensor networks: Fundamental concepts and state of the art. *ACM Computing Surveys (CSUR)*, 43(3):1–51, 2011. 2
 - [19] Ryan Newton, Greg Morrisett, and Matt Welsh. The regiment macroprogramming system. In *2007 6th International Symposium on Information Processing in Sensor Networks*, pages 489–498. IEEE, 2007. 2
 - [20] Rishiyur S Nikhil et al. Executing a program on the mit tagged-token dataflow architecture. *IEEE Transactions on computers*, 39(3):300–318, 1990. 1, 2, 3
 - [21] Kamin Whitehouse, Cory Sharp, Eric Brewer, and David Culler. Hood: a neighborhood abstraction for sensor networks. In *Proceedings of the 2nd international conference on Mobile systems, applications, and services*, pages 99–110, 2004. 2
 - [22] Jennifer Yick, Biswanath Mukherjee, and Dipak Ghosal. Wireless sensor network survey. *Computer networks*, 52(12):2292–2330, 2008. 2
 - [23] Jia Zou, Slobodan Matic, Edward A Lee, Thomas Huining Feng, and Patricia Derler. Execution strategies for ptides, a programming model for distributed embedded systems. In *2009 15th IEEE Real-Time and Embedded Technology and Applications Symposium*, pages 77–86. IEEE, 2009. 2